

ILLVZN Final Pitch

HackRx 6.0



2025

~ Avinash Anish & Nihaal SP (RVCE)



Problem Statement

★ The Real Challenge

Build a system that uses Large Language Models (LLMs) to process natural language queries and retrieve relevant information from large unstructured documents

The Hidden Crisis

- Information Asymmetry
- Context Collapse
- Human Bottlenecks
- Trust Deficit



Round 1, RAG?

- Entered hackathon with the idea that It would be a very fun problem statement to solve, and boy it was.
- Limited understanding of retrieval, chunking, and embeddings didn't help initially.
- Used Pinecone for vector storage → faced latency & bottlenecks.
- Relied on a local embedding model with poor accuracy/quality.
- Outcome: Only 1 low scoring performance in Round 1 which almost buried us
- R1 accuracy: 4% and a rank of 323:(

Round 2, Document Bombardment

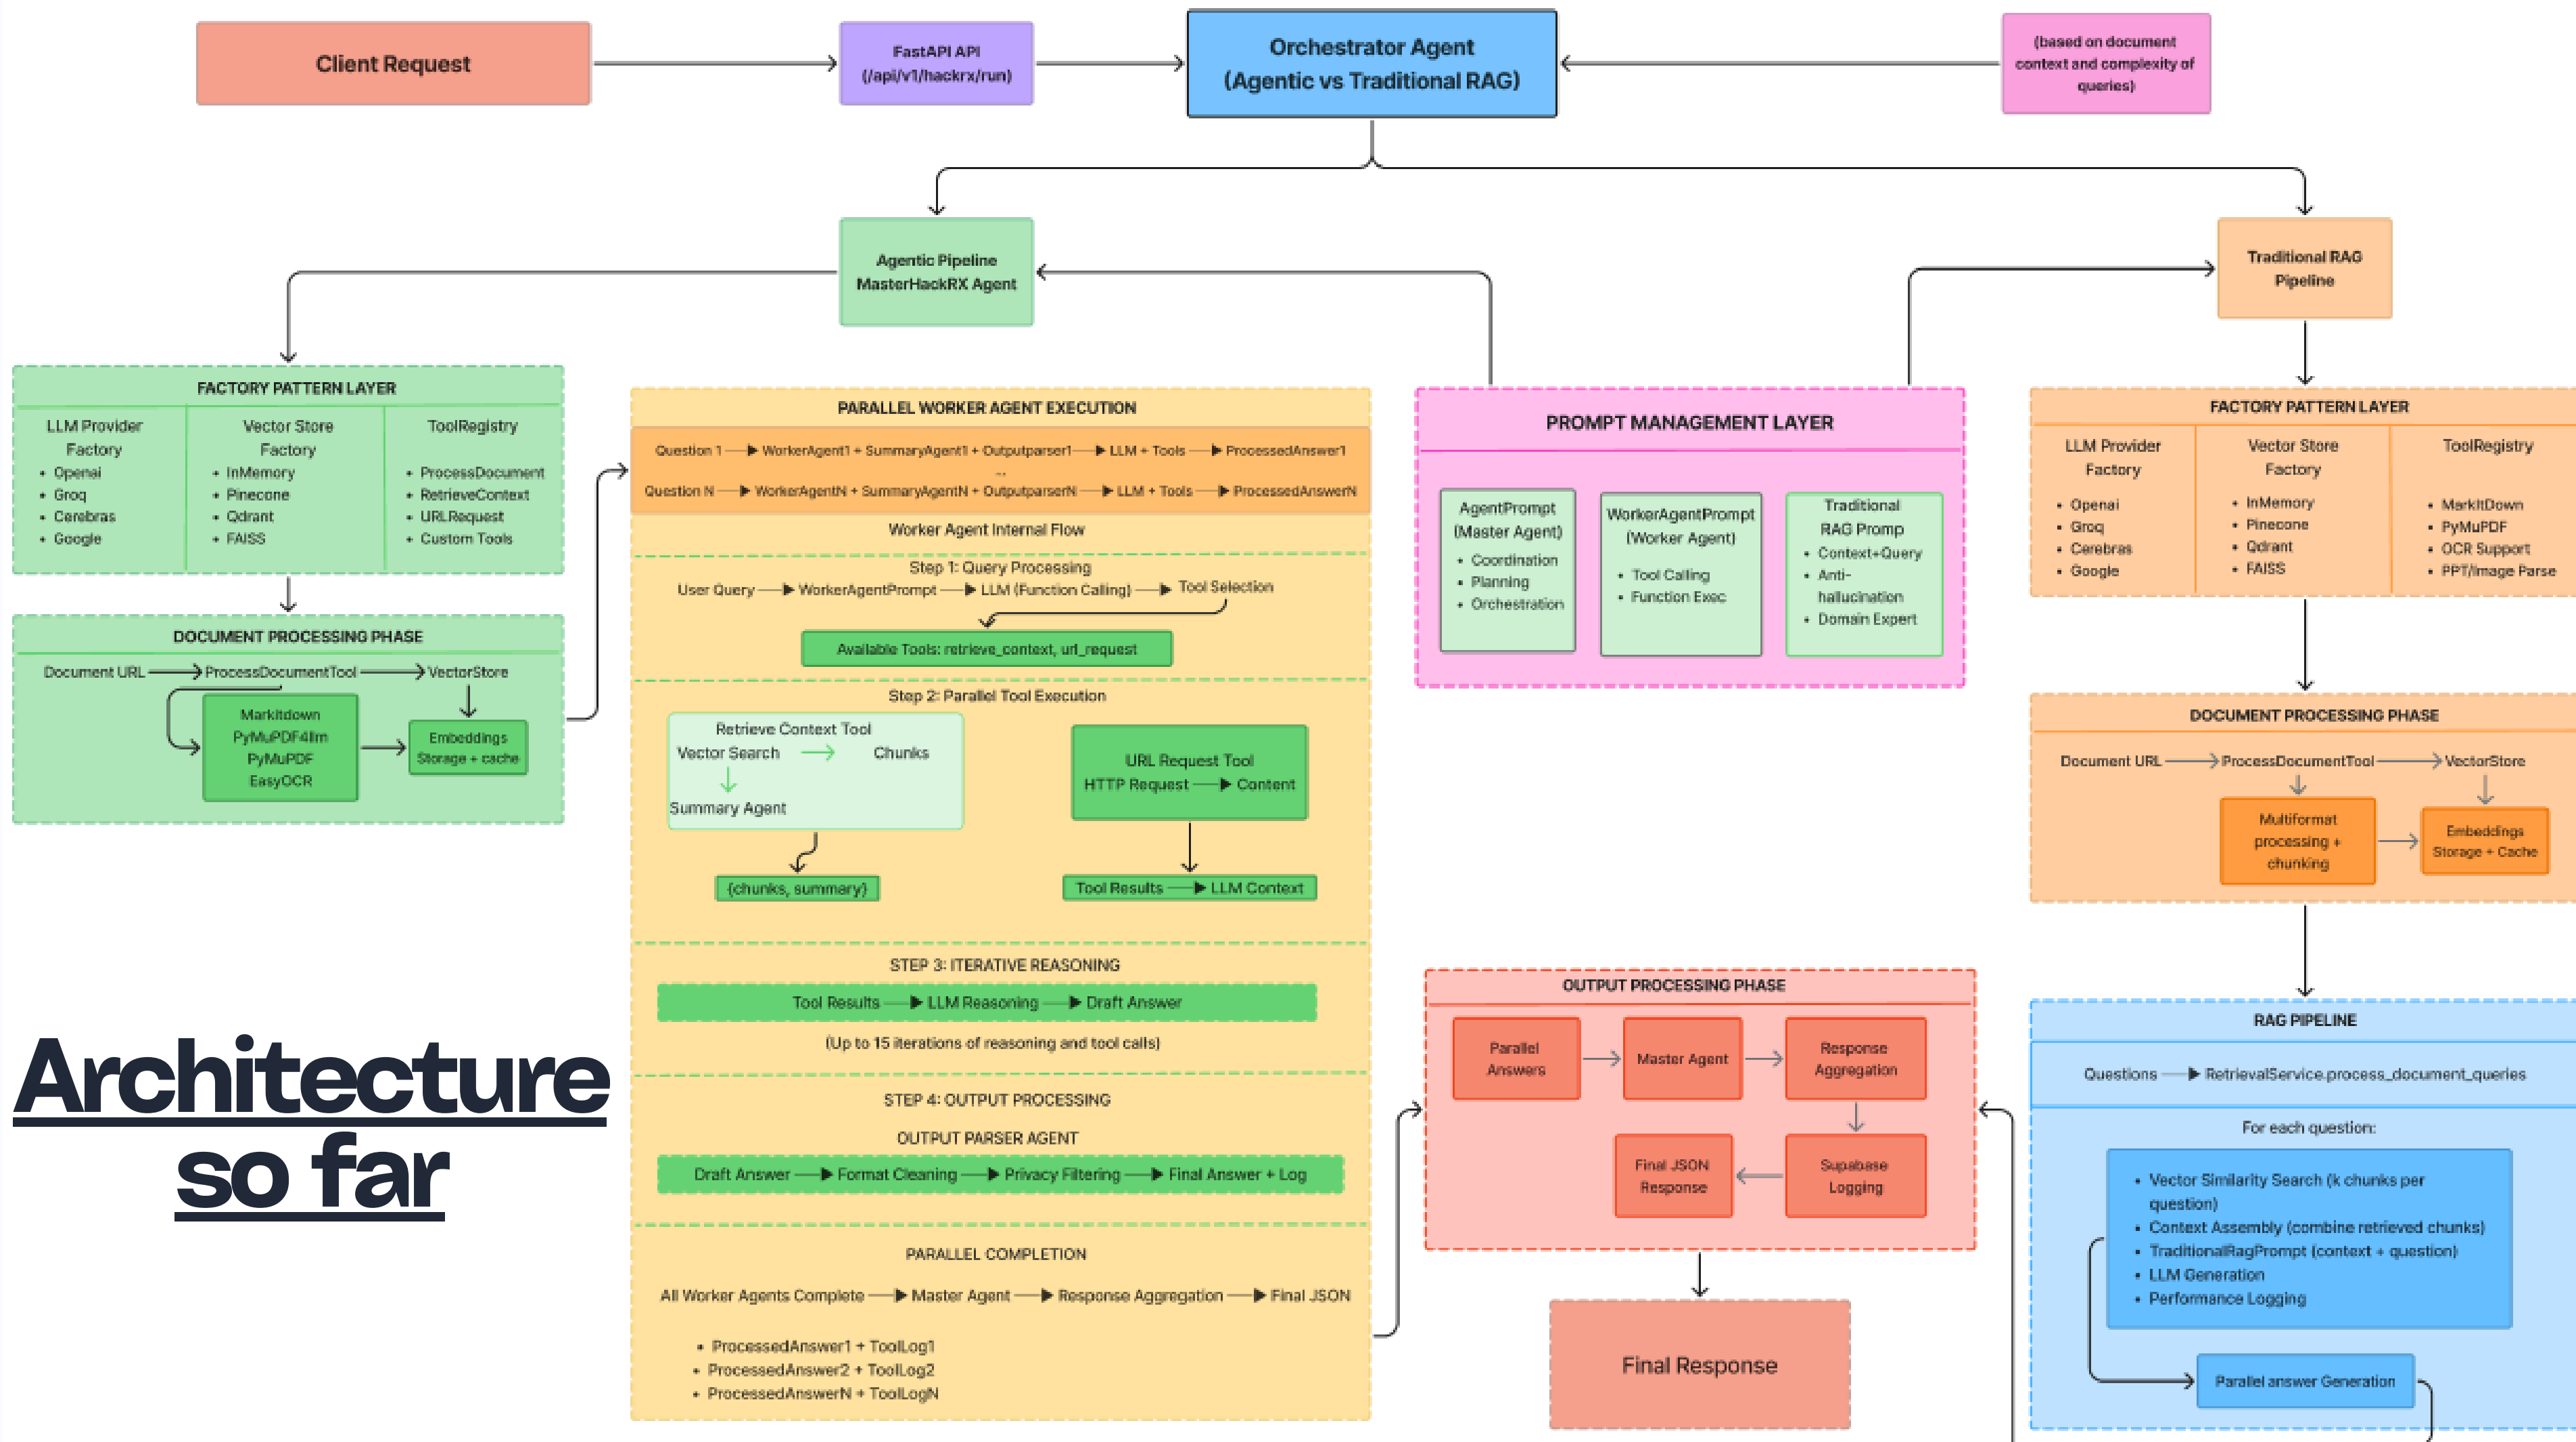
- Now having a good idea about how things worked, both we and the number of documents entered into a higher gear.
- Switching to inmemory vector DB, Openai embedding models, and implementing async parallelisation really worked in our favour
- We topped the leader board in accuracy, and were 5th score wise and that's when we realised our standards and scope of HackRx
- At this point, we also started logging the questions and responses onto Supabase for debugging any issues and made testing scripts which can hit our endpoint in the same way.
- R2 accuracy: 66% and 1st rank(by accuracy), 5th(by score)

Round 3, the Rain continues


- Round 2 had some huge documents which we were just able to abide by, but then Round 3 added onto that with all kinds of documents like PPTs, Images, binary dumps, etc
- Added factory patterns for LLM providers, vector DBs, better document processing libraries and any small optimisations we could think of
- All our scores were close, and we probably should have thought about using better opensource embedding models which we now realise perform much better than the Openai embeddings.
- R3 accuracy: 75%

Round 4, the Enigma

- Until round 3 all queries were based off of documents that our Traditional RAG system could solve, but the queries of Round 4 were on a different level.
- We then realised our naivety and the small reach of our architecture and decided we needed to improve as R4 needed agentic tool calls.
- Then things fell in correctly, and we added an agentic pipeline along with the traditional RAG system with a master orchestrator to make the choice for the pipeline to be used based on the query and its complexity
- A common prompt layer provided prompts for both the pipelines to encourage Generalism, and the agentic pipeline had worker agents performing the tool calls
- With agentic comes latency which was our bane this round and we almost didn't make it to the final showdown.
- R4 accuracy: 84.31%



Tech Stack



NEXT.js



Gemini



Round 5, Waste of potential

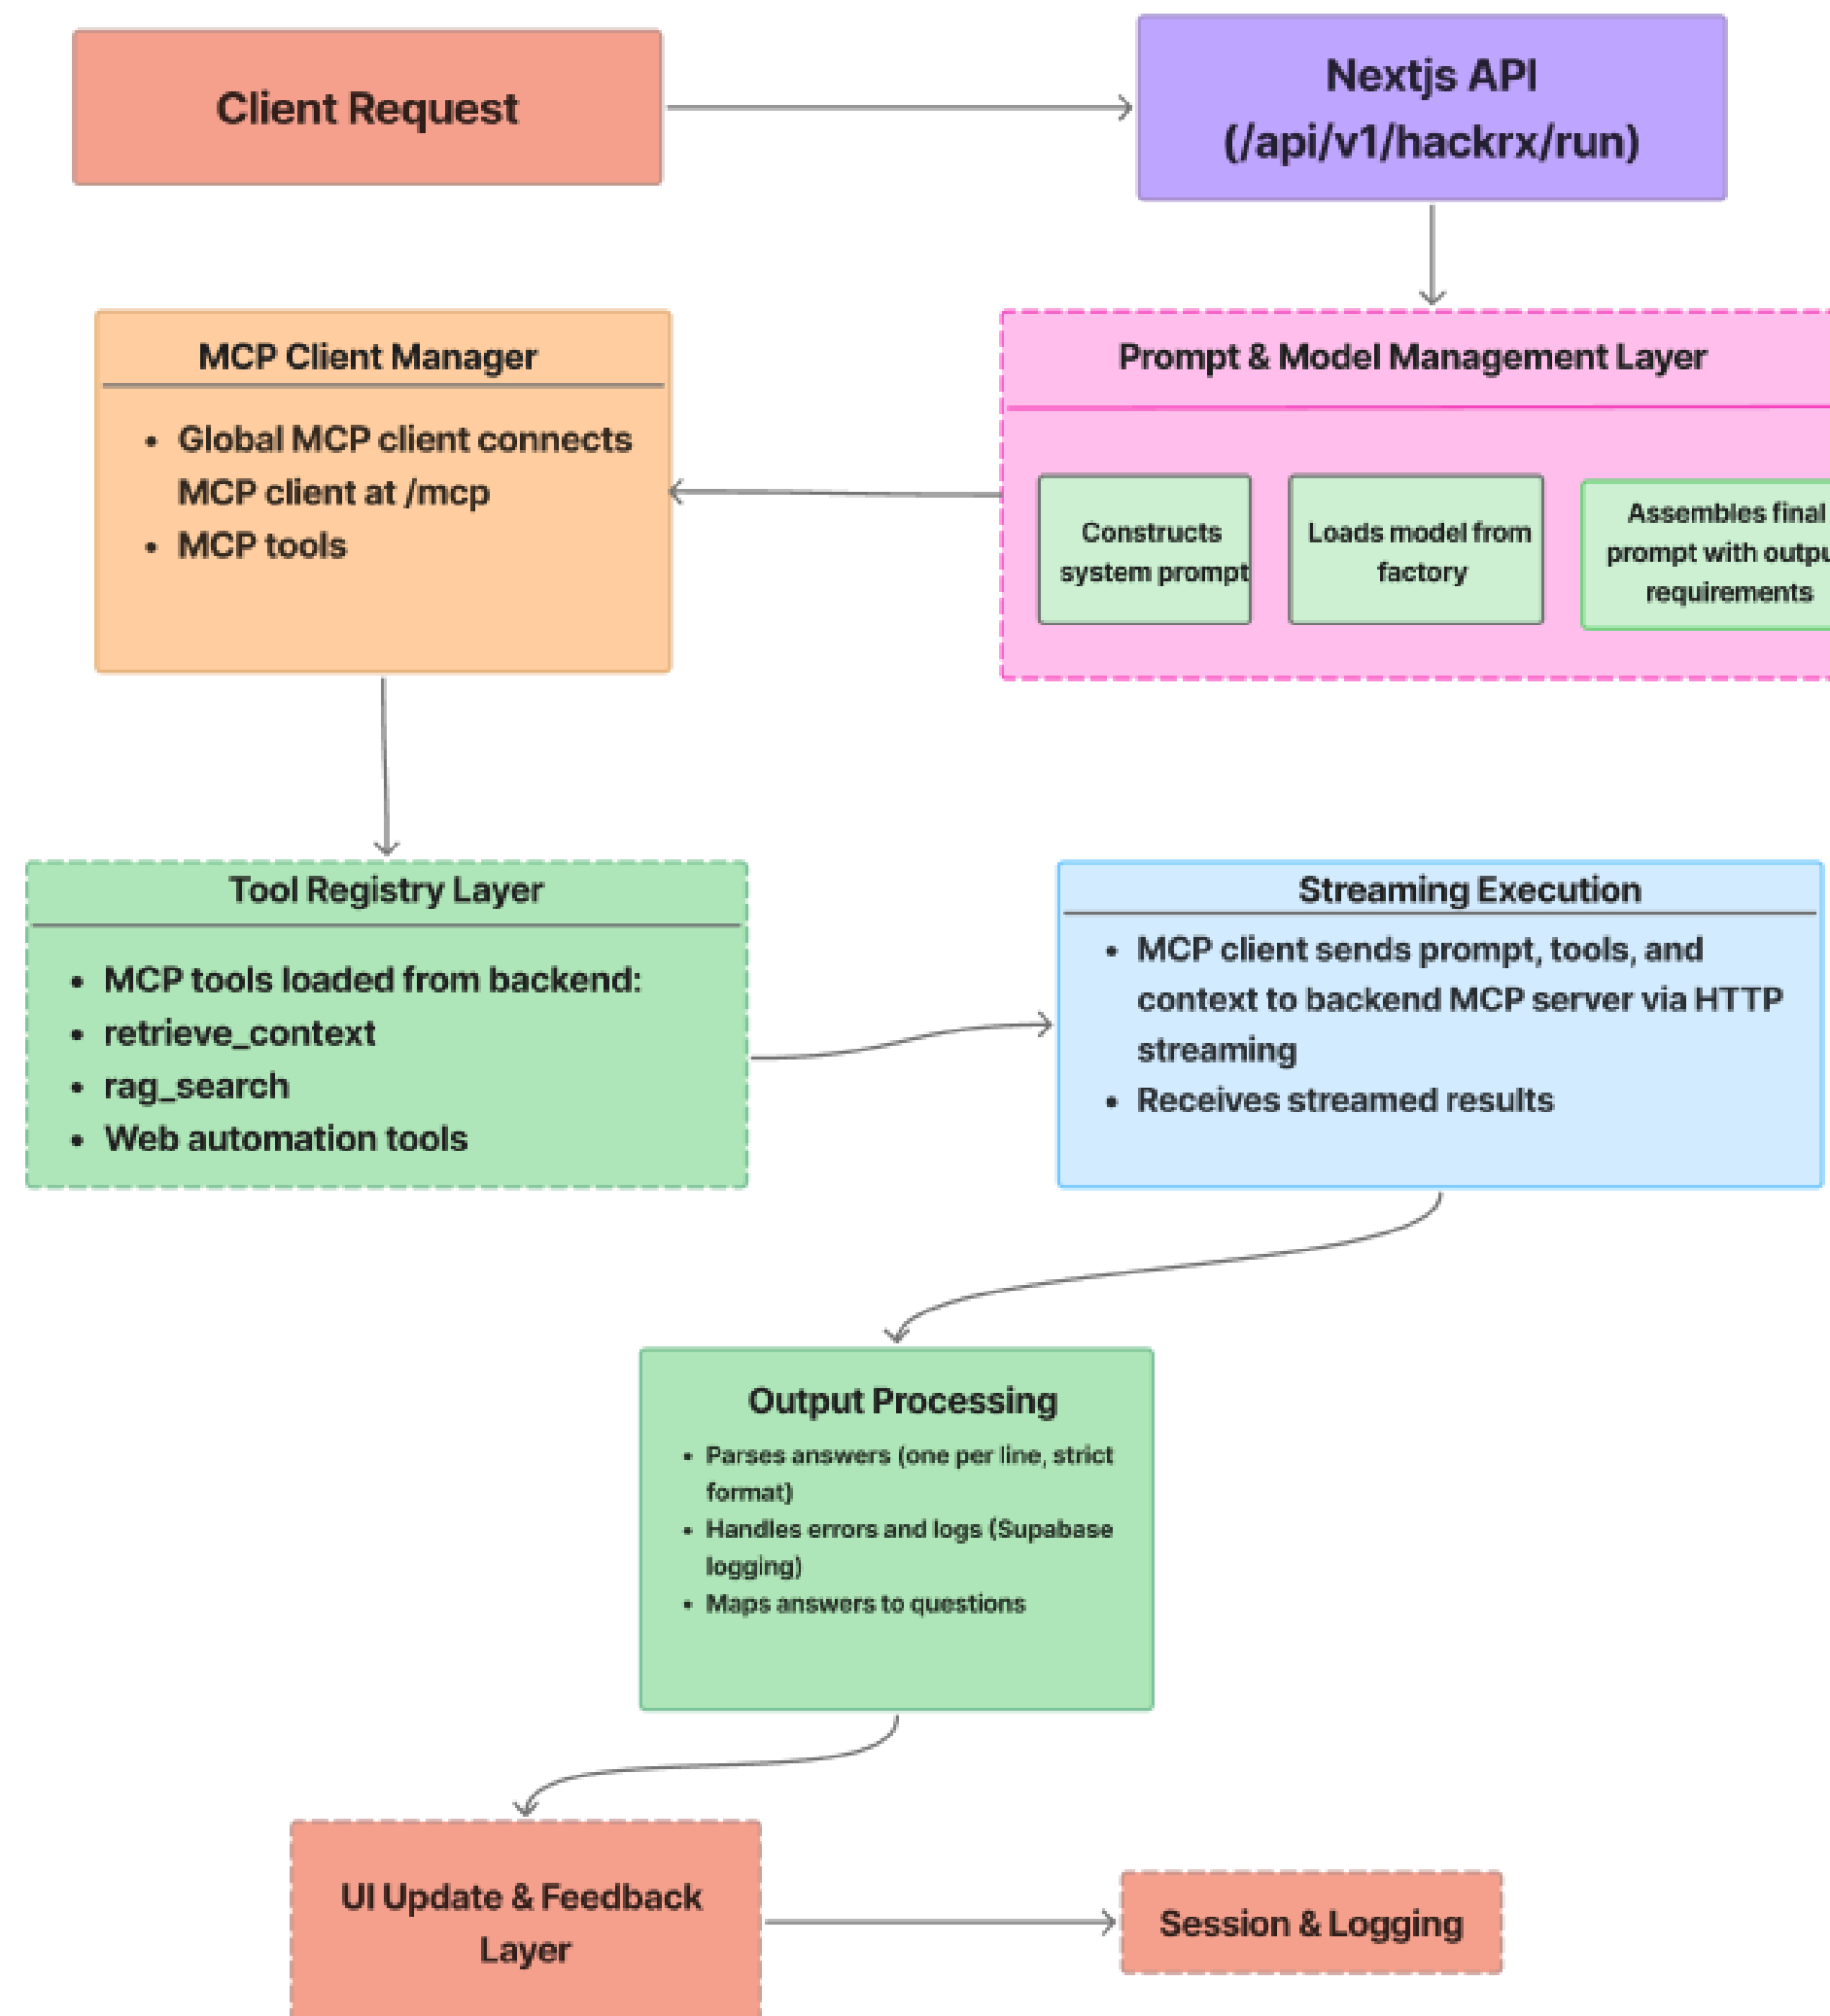
- Round 5 kicked off with a token based game where we had to find the hidden text in the source code.
- Since the code was client side, we generated valid JWTs which opened 3 more games to us which were actually meant to be R7 questions lol
- Instead of submitting, we kept testing to no avail as the server was decoding our token, due to which the game couldn't be played and we wasted hours
- When we made our first submission, it immediately succeeded which left us in mixed emotions of regret for wasting time but also to strive even harder and not make such rookie mistakes again
- Initialised Playwright MCP for all browser based actions
- Achieved blazing fast speeds but were moved onto the penultimate round by a margin of 0.42
- R5 accuracy: 85%

Round 6, Agentic Dev

- Wrote more MCP tools, like a code execution and a Github commit and push tools.
- The agent had to solve a question, run it with test cases, return the answer and push it to our private HackRx repos.
- As usual everything can't go as planned, and we forget to verify the response format, and hadn't catered for the URL field(it was an empty string) and were hit with a score of 0 after 0 due to the Azure content policy warnings.
- After several failed attempts, it clicked to us at 3am and we struck gold with a successful submission and off to round 7 we went.
- R6 accuracy: 95%

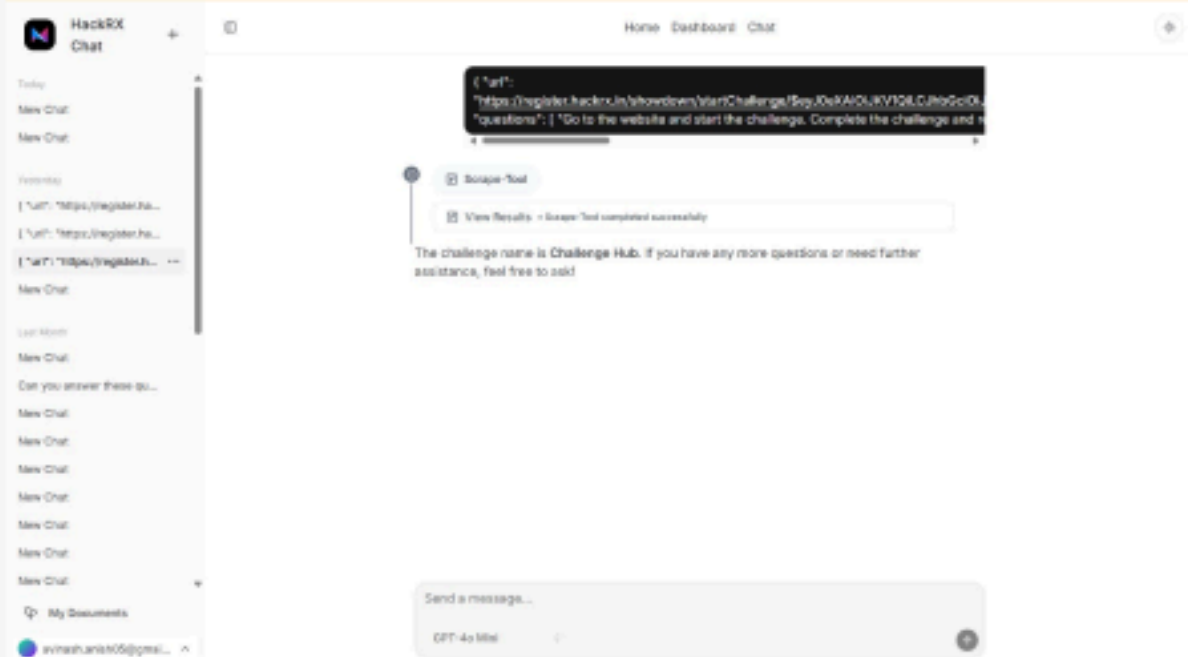
Round 7, Final Push

- Despite knowing the questions from the start, we simply couldn't solve the Simon says pattern matching question, but solved the Github repository analysis question.
- Probably could've written a tool to check the CSS mutations to solve it but felt it would be too specific and would go against the rules of the hackathon to make everything generalised. Thus, we finished 10th on the leader board.
- The mentorship was phenomenal, and no matter the result ultimately at the end of the day we learnt a lot, and will definitely build some agentic pipelines with MCP integrations in the future.
- R7 accuracy: 79.17%



Architecture
at the end of
Round7

Tech Stack- Frontend



TS TypeScript

Tailwind CSS

React

NEXT.js

Database



supabase

Database

Tech Stack- Backend



FastAPI



PyTorch



LangChain



OpenAI

Deployment

ngrok



Vertex AI



Render

Technology
Stack

Some Metrics

Average tokens per Agentic Task - GPT 4.1-mini/4o-mini



Round 5 - 50k in, 703 out
Average Time: 28 s

» Input

Query: Go to the website and start the challenge. Complete the challenge and return the answers for the following questions. Questions: 1. What is

Show More

« Output

completed

sequence

88b189

LLM Usage Statistics:

Prompt	Completion	Total
50703 tokens	703 tokens	51406 tokens

HackRXAgent

System Prompt Runtime

You are HackRXAgent. You are an AI agent specialized in web automation, code...

+ Details

Status: completed Completed

Model: gpt-4o-mini



Round 6 - 18k in, 812 out
Average Time: 8 s

» Input

Query: You are given two strings s and pattern. A string x is called almost equal to y if you can change at most one character in x to make it

Show More

« Output

completed

1

4

LLM Usage Statistics:

Prompt	Completion	Total
18198 tokens	812 tokens	19010 tokens

HackRXAgent

System Prompt Runtime

You are HackRXAgent. You are an AI agent specialized in web automation, code...

+ Details

Status: completed Completed

Model: gpt-4o-mini



Round 7 - 50k in, 705 out
Average Time: 30s

» Input

Query: Go to the website and start the challenge. Complete the challenge and return the answers for the following questions. Questions: 1. What is

Show More

« Output

completed

sequence

6ee3e9

LLM Usage Statistics:

Prompt	Completion	Total
50969 tokens	705 tokens	51674 tokens

HackRXAgent

System Prompt Runtime

You are HackRXAgent. You are an AI agent specialized in web automation, code...

+ Details

Status: completed Completed

Model: gpt-4o-mini

Debug Info: ID: 6d02cac8-ca45-4b33-9ba6-618a3cf355b4 Status: completed Event: 24/24 {...} Run


Our USP

How we stand out



Multi Agent System

The combination of orchestrator, query enhancers, output parsers and worker agents makes our system fully autonomous due to the



MCP Server Integration

We integrate MCP Servers from various providers along with custom tools, enabling our agentic system to perform complex tasks with tool calling




Observability

Integration of the TypeScript Multi Agent framework - Volt Agent, allows us to have unmatched observability, with request level logging and orchestration visualisations



Parallel Execution

Parallel execution of user queries ensures that our system provides answers to the user's queries in less than 20 seconds on average without streaming



Scalability Aspects

1. Horizontal Architecture Scalability

- Container-Ready Deployment
- Stateless API Design

2. Vector Store & Data Scalability

- Multi-Factory VectorDB Support
- Intelligent Caching Strategy

3. AI Model & Provider Scalability

- Multi-Provider Architecture
- Provider Factory Pattern

4. Auto-Scaling & Load Management

- Kubernetes Ready
- Health Check Endpoints
- Graceful shutdown

Thank You

